


LLM Fundamentals: Under the Hood of Generative Language Models

Modern LLMs are not databases or search engines: they are statistical token predictors governed by context and attention.



TAKEAWAY: An LLM does not retrieve stored facts: it computes probability distributions over tokens conditioned on input context.


Fluency Is Not Knowledge: The Missing Evidence Trap

When asked to summarize an unprovided policy, an LLM outputs coherent, grammatically flawless, yet fully fabricated text.

ANTI-PATTERN: UNGROUNDED INFERENCE



ENGINEERING PATTERN: GROUNDED SYNTHESIS




TAKEAWAY: Fluency is cheap; grounding is architectural. Without context, generation is probabilistic guesswork.

The Autoregressive Loop: Sequential Token Sampling

Generation is strictly sequential: each new token is appended to the input buffer to compute the subsequent token probability.

Autoregressive Step Rollout

Sequential accumulation in context buffer



Latency Breakdown

Compute-bound vs Memory-bound regimes

Phase 1: Prefill (TTFT)

Parallel forward pass over all N prompt tokens.
Compute-Bound (Tensor Cores 100%)

Phase 2: Decoding (ITL)

Sequential 1-token passes. Generates 1 token per cycle.
Memory-Bandwidth Bound (VRAM read per token)
Inter-Token Latency (ITL) ~ 20ms - 40ms / token

Engineering Latency Formula:

$Latency_{total} = TTFT + (N_{generated} * ITL)$

TAKEAWAY: Inference latency is proportional to output length: Time-to-First-Token (prefill) differs fundamentally from Inter-Token Latency (decode).

Why Subwords: The Vocabulary-Coverage Tradeoff


Character-level models create prohibitive sequence lengths; word-level models explode in vocabulary size. Subwords balance both.

1. Character-Level	2. Word-Level	3. Subword (BPE / WordPiece)
Vocabulary Size: ~256 bytes (Extremely compact) Sequence Length (L): 4x - 6x longer sequences "transformer" = 11 characters Computational Penalty: O(L^2) Attention explodes 16x - 36x compute and KV-cache Bottleneck Weak semantic signal per step. Models struggle with long-range reasoning.	Vocabulary Size: > 1,000,000 words (Massive) Sequence Length (L): Shortest sequence length 1 word = 1 token Embedding Memory: Embedding matrix gigabytes Weights dominated by sparse lookup Bottleneck Out-of-Vocabulary (OOV) disaster. Any typo or new code symbol maps to <UNK>.	Vocabulary Size: 32k - 128k tokens (Bounded) Sequence Length (L): Near-word compression ratio English ~1.3 tokens per word Coverage Property: Zero Out-of-Vocabulary (0% OOV) Decomposes unseen words into characters Industry Standard Optimal efficiency and semantic density. Powers GPT-4, Llama 3, Claude, Gemini.

TAKEAWAY: Subword tokenization strikes the optimal balance between sequence length and vocabulary memory footprints.

Byte-Pair Encoding: From Character Chunks to Token IDs

BPE iteratively merges the most frequent adjacent character pairs until reaching target vocabulary capacity.



TAKEAWAY: Words are segmented into statistical chunks: rare words fragment into multiple sub-tokens, increasing cost and latency.

Token Economics: Throughput, Cost, and Modality Skew

Tokens govern GPU memory bandwidth, cloud API billing, and latency: non-Latin scripts pay an implicit 2x-4x token penalty.

Token Consumption for Identical Semantic Content

Sentence: "System optimization reduces server latency."

English (Baseline) 5 Tokens (1.0x)

5 Tokens ["System", " optimization", " reduces", " server", " latency."]

Cyrillic Script (Russian) 14 Tokens (2.8x Tax)

14 Tokens (Fragmented sub-chunks)

Fragments: ["Оп", "timi", "зация", " си", "стемы", " сни", "жает", ...]

Python Code with Indentation 9 Tokens (1.8x Tax)

9 Tokens

Separate tokens spent on leading whitespace and symbol delimiters

Impact: 2.8x token count = 2.8x higher API cost + 2.8x longer decode latency

API Unit Economics

Standard cloud provider cost model

Input vs Output Pricing

Input: \$0.15 - \$2.50 / 1M tokens
Output: \$0.60 - \$10.00 / 1M tokens (4x)

Throughput & Rate Limits

TPM: Tokens Per Minute quota
RPM: Requests Per Minute quota
Exceeding TPM throws HTTP 429 Rate Limit


Engineering Sizing Rule

Calculate token budget before injecting documents:
 $Max_Docs = Context_Budget - Output_Reserve$

TAKEAWAY: Token efficiency varies drastically by script and formatting: system budgets must account for language asymmetry.

From Discrete IDs to Continuous Vectors: Input Embeddings & RoPE


Token IDs are mapped to high-dimensional vectors, enriched with rotary positional encodings before entering attention layers.



TAKEAWAY: Embeddings translate discrete symbols into continuous geometry: positional rotations preserve grammatical sequence order.


Self-Attention Mechanics: Dynamic Routing of Information

Attention computes how much each token focuses on every other token via Query, Key, and Value dot-product matrices.



Multi-Head Attention: Attending to Multiple Relations in Parallel

Splitting hidden dimension d across h attention heads allows the model to simultaneously track syntax, co-reference, and factual relations.



TAKEAWAY: Multiple attention heads project tokens into distinct representational subspaces in parallel, merging through an output projection matrix.

Quadratic Complexity: The Memory and Compute Wall

Standard self-attention scales quadratically with sequence length: long context requires memory-efficient attention algorithms.

The $O(L^2)$ Quadratic Wall

Matrix elements and VRAM scaling per layer

$L = 2,048$ Tokens (Early LLMs)

Matrix Size: $2k \times 2k = 4.19M$ entries

~16 MB / layer

$L = 32,768$ Tokens (Modern Baseline)

Matrix Size: $32k \times 32k = 1.07B$ entries

~4.2 GB / layer

$L = 131,072$ Tokens (Long Context)

Matrix Size: $128k \times 128k = 17.18B$ entries

~68.7 GB / layer

In 80-layer models, standard attention crashes VRAM (OOM Error)

Architectural Solution: FlashAttention

IO-Aware Kernel: Eliminating HBM round-trips

GPU High Bandwidth Memory (HBM, 80 GB)

Large capacity, but slow bandwidth bottleneck (~3 TB/s)



Standard: Repeated accesses: 1 read, 1 write



On-Chip SRAM Cache (Ultra Fast, ~19 TB/s)

Matrix split into small tiles (e.g. 128×128)

Online Softmax: Tiling avoids full $L \times L$ matrix materialization


Efficiency Result:

Memory drops from $O(L^2)$ to $O(L)$ | 3x to 5x speedup

TAKEAWAY: Attention computation and memory scale as $O(L^2)$: FlashAttention bypasses memory bandwidth walls by tiling inside GPU SRAM.


The KV-Cache: Eliminating Redundant Attention Computation

Prefill computes keys and values for all prompt tokens once: decoding reuses stored vectors from GPU VRAM instead of recalculating.



Anatomy of the Context Window: Allocating Working Memory

The context window is a finite budget shared between system directives, retrieved evidence, history, and output room.



TAKEAWAY: The context window is working memory, not a data lake: saturating context with irrelevant tokens degrades reasoning precision.

In-Context Conditioning vs Parameter Weights

Adding text to a prompt modifies activations for a single forward pass: it does not update weights or retain persistent state.

IN-CONTEXT CONDITIONING

Ephemeral Working Memory

Dynamic runtime activations inside attention layers

Persistence:
Zero Persistence (Volatile in VRAM, vanishes on request end)

Weight Update:
 $\Delta W = 0$ (Strictly frozen parameters, no backpropagation)

Mechanism:
Attention mechanism steers next-token probabilities via prompt K, V

Analogy: CPU Registers / RAM Scratchpad
Data is active only while the application process runs

PARAMETRIC WEIGHTS

Permanent Frozen Checkpoint

Static matrix weights stored in model safetensors

Persistence:
Permanent (Persisted on disk across millions of sessions)

Update Cost:
Massive (Pre-training or Fine-Tuning with cluster GPUs)


Mechanism:
Matrix multipliers baked during gradient descent optimization

Analogy: Hardwired ROM / Compiled Binary
Immutable during inference execution

TAKEAWAY: Prompts never modify model weights. What is absent from context or trained weights does not exist for the model.

Context Degradation: The "Lost in the Middle" Phenomenon

Empirical studies reveal a sharp U-curve: models retrieve facts at the start and end of context far better than in the middle.



Boundary Engineering Rules

How to structure long prompt contexts

1. Pin Rules to Start (Primacy)

Core instructions, role, and output schemas
Primacy: Highest attention at prompt start

2. Neutral Middle Chunk Space

Reference docs & auxiliary text chunks here
Caution: Avoid critical logic in center

3. Pin Task to End (Recency)

Repeat exact task instructions and schemas
Recency: Direct proximity to generation

Reranking: Put high-score chunks at boundaries

TAKEAWAY: Position critical instructions and primary evidence at prompt boundaries: avoid burying essential facts in the center.

Prompt Anatomy: Precision Task Specification

High-reliability prompts are structured specifications with clear role, objective, evidence, explicit constraints, and verification schemas.

MODULE 1: ROLE

System Persona & Tone

Sets vocabulary priors and domain boundaries

```
"You are an expert reliability engineer"
"adhering strictly to ISO-26262 safety."
```

MODULE 2: CONTEXT

Delimited Source Data

Isolates untrusted data from instructions

```
<context id="incident_report">
Timestamp: 2026-09-20T14:12Z ...
</context>
```

MODULE 3: TASK

Deterministic Objective

Verifiable transformation or extraction

```
"Extract primary root-cause code"
"and calculate outage duration in ms."
```

MODULE 4: CONSTRAINTS

Negative Boundaries (Do NOT)

Eliminates assumptions and default filler

```
"Do NOT invent missing numbers."
"Output NULL if metric is absent."
"Zero introductory text or conversational p
```

MODULE 5: EXEMPLARS

Few-Shot Golden Pairs

Demonstrates edge case handling directly

```
Input: "Disk full on /var/log"
Output: {"code": "ERR_FS_FULL", "sev": 1}
Pattern anchors formatting style
```

MODULE 6: SCHEMA


Strict Output Validation

Pydantic / JSON Schema contract

```
{"type": "object",
"required": ["root_cause", "duration_ms"]}
Guarantees deterministic JSON parsing
```

The Prompting Strategy Tree: Matching Complexity to Technique

Different operational workloads require different reasoning scaffolds: from zero-shot queries to reasoning traces and ReAct tool loops.



TAKEAWAY: Escalate prompting complexity intentionally: add few-shot for formatting, CoT for multi-step logic, and ReAct for tool execution.

Few-Shot Calibration: The Balance Between Guidance and Bias

Examples guide structure effectively, but imbalanced classes, formatting noise, or ordering artifacts bias model classification.

CALIBRATED EXEMPLAR SET

Balanced Statistical Prior

Equal class distribution across target taxonomy

- Ex 1 [BUG]: "Null pointer in auth service" → BUG
- Ex 2 [FEATURE]: "Support OAuth2 PKCE flow" → FEATURE
- Ex 3 [BILLING]: "Invoice receipt missing VAT" → BILLING

Model Classification Accuracy: 94.2%
Unbiased probability distribution across categories.
 $P(\text{Bug}) = 0.33, P(\text{Feature}) = 0.33, P(\text{Billing}) = 0.33$

SKEWED EXEMPLAR SET

Unintended Prior Skew

Label imbalance trains the model to over-predict bugs

- Ex 1 [BUG]: "Null pointer in auth service" → BUG
- Ex 2 [BUG]: "Database connection timed out" → BUG
- Ex 3 [BUG]: "Crash on image upload resize" → BUG

Severe Misclassification: Billing queries labeled as BUG
Model internal prior heavily skewed by majority class.
 $P(\text{Bug}) = 0.91, P(\text{Feature}) = 0.05, P(\text{Billing}) = 0.04$


TAKEAWAY: Few-shot examples act as statistical priors: ensure class balance, uniform length, and clean label distributions to avoid skew.

Vector Embeddings: Mapping Semantic Distance to Geometry

Embedding models map text strings into high-dimensional vector space where conceptual similarity equals geometric proximity.

Vector Space Projection (Semantic Clustering)

High-dimensional embedding vectors projected onto 2D plane



Cosine Similarity Metric

Normalized dot product measuring directional alignment

$$\text{Cosine_Sim}(u, v) = (u \cdot v) / (||u|| * ||v||)$$

Measures orientation, strictly invariant to vector magnitude

Score Interpretation Spectrum:

- +1.00 : Identical semantic direction
- 0.00 : Orthogonal (Completely unrelated topics)
- 1.00 : Diametrically opposite meaning


Operational Principle

Embeddings do NOT write text: they retrieve candidate evidence.

TAKEAWAY: Embeddings convert linguistic meaning into spatial coordinates: cosine similarity measures conceptual proximity.

Lexical vs Semantic Retrieval: Why Production Needs Hybrid

Dense vector search excels at conceptual matching but misses exact alphanumeric identifiers: BM25 lexical search provides the inverse.




Vector Storage & Indexing: Scaling to Millions of Vectors

Exhaustive brute-force cosine distance is intractable at scale: HNSW graph indexes enable sub-millisecond retrieval across millions of vectors.

HNSW: Multi-Layer Skip Graph Architecture

Fast coarse navigation transitioning to dense neighborhood search



Algorithm Tradeoffs

Exact Scan vs HNSW Indexing

Flat Index (Exact Brute-Force)

Time Complexity: $O(N * d)$
Recall: 100% | Latency at 10M vectors: ~850ms

HNSW Index (Approximate Graph)

Time Complexity: $O(\log N * d)$
Recall: 98.5% | Latency at 10M vectors: ~4ms
200x query speedup with <1.5% recall loss


Tunable Engineering Knobs:

M: Number of bidirectional links per node (16-64)
efConstruction: Build search depth (100-400)
efSearch: Runtime query beam width (32-128)

TAKEAWAY: Vector databases trade marginal recall precision for logarithmic search latency using graph-based ANN indexes like HNSW.

The End-to-End RAG Pipeline: Ingestion and Query Execution


RAG separates static reasoning from dynamic knowledge: documents are chunked and indexed offline, then retrieved dynamically at query time.



TAKEAWAY: RAG grounds generation in verifiable enterprise data: the database provides truth, the model provides synthesis and language.

Failure Modes Across the RAG Pipeline: Where Systems Break


Production RAG failures rarely originate in model generation: they stem from chunking boundary loss, stale indices, and retrieval misses.



TAKEAWAY: Most RAG bugs are retrieval and chunking bugs. Grounding is only as reliable as the retrieved context chunks.

Context Security: Mitigating Direct and Indirect Prompt Injection

Untrusted text retrieved into context can hijack system directives: rigorous boundary separation and post-generation scanners are mandatory.



TAKEAWAY: Treat external retrieved context as untrusted user input: isolate data with strict XML delimiters and output guardrails.

The LLM Engineering Matrix: Prompting, RAG, Fine-Tuning, or Agents

Select the minimal architectural complexity required to satisfy domain requirements, latency budgets, and accuracy thresholds.

TIER 1: PROMPT ENGINEERING

Zero Infrastructure Baseline

Direct prompting, few-shot exemplars, structured JSON

Setup Cost: Zero | Latency: 1 Call | Agility: Fast
Best for: Fast POC, summarization, text transforms
Constraint: Bound by context and training cutoff

TIER 2: GROUNDED RAG

Dynamic Knowledge Ingress

Semantic chunking + hybrid retrieval (Vector + BM25)

Setup Cost: Medium | Latency: +50-150ms | Citations: 100%
Best for: Knowledge bases, compliance, audit
Decouples knowledge from static neural weights

TIER 3: FINE-TUNING

Style & Grammar Specialization (LoRA)

Adapting model weights to specialized syntax and concise output

Setup Cost: High (GPU) | Latency: Lowest | Format: Exact
Best for: Custom syntax, compact output, edge models
Notice: Does NOT inject facts; use RAG for recall

TIER 4: AUTONOMOUS AGENTS

Tool Execution & State Loops

ReAct cycles, function calling, external database modifications

Setup Cost: Highest | Latency: Multi-step | Autonomy: Max
Best for: Multi-API tasks, triage, tool loops
Requirement: Human-in-the-loop safety fences

TAKEAWAY: Never deploy complex architecture when a simpler primitive suffices. Start with prompting, scale to RAG, fine-tune for latency.