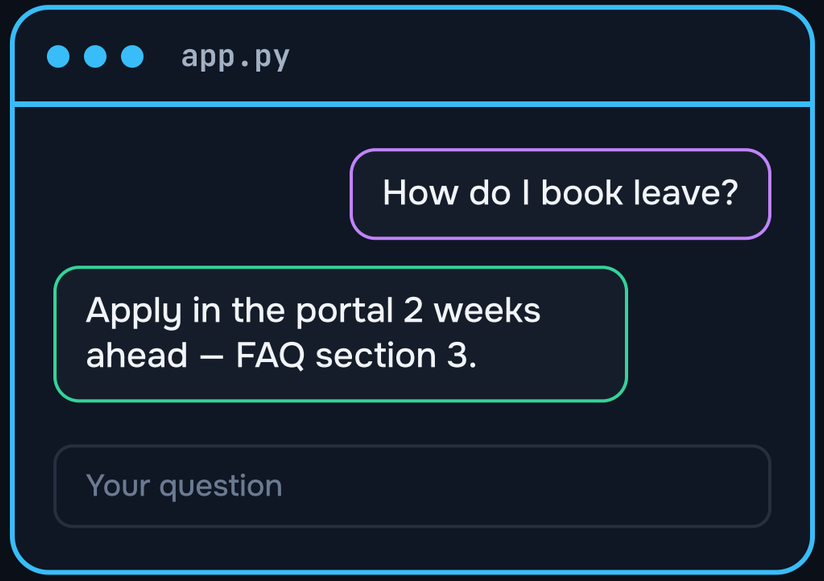
GENAI BASICS • LESSON 14 • MODULE 6: IDEA TO PRODUCT

Rapid AI app prototype

Your own chat window with Streamlit in a day — and when a no-code builder is enough

Streamlit, a coding agent, publishing, secrets. Google AI Studio Build, Claude Artifacts. Data as of October 2026.



WHY

A prototype answers “what will it look like?” in a day, not a quarter

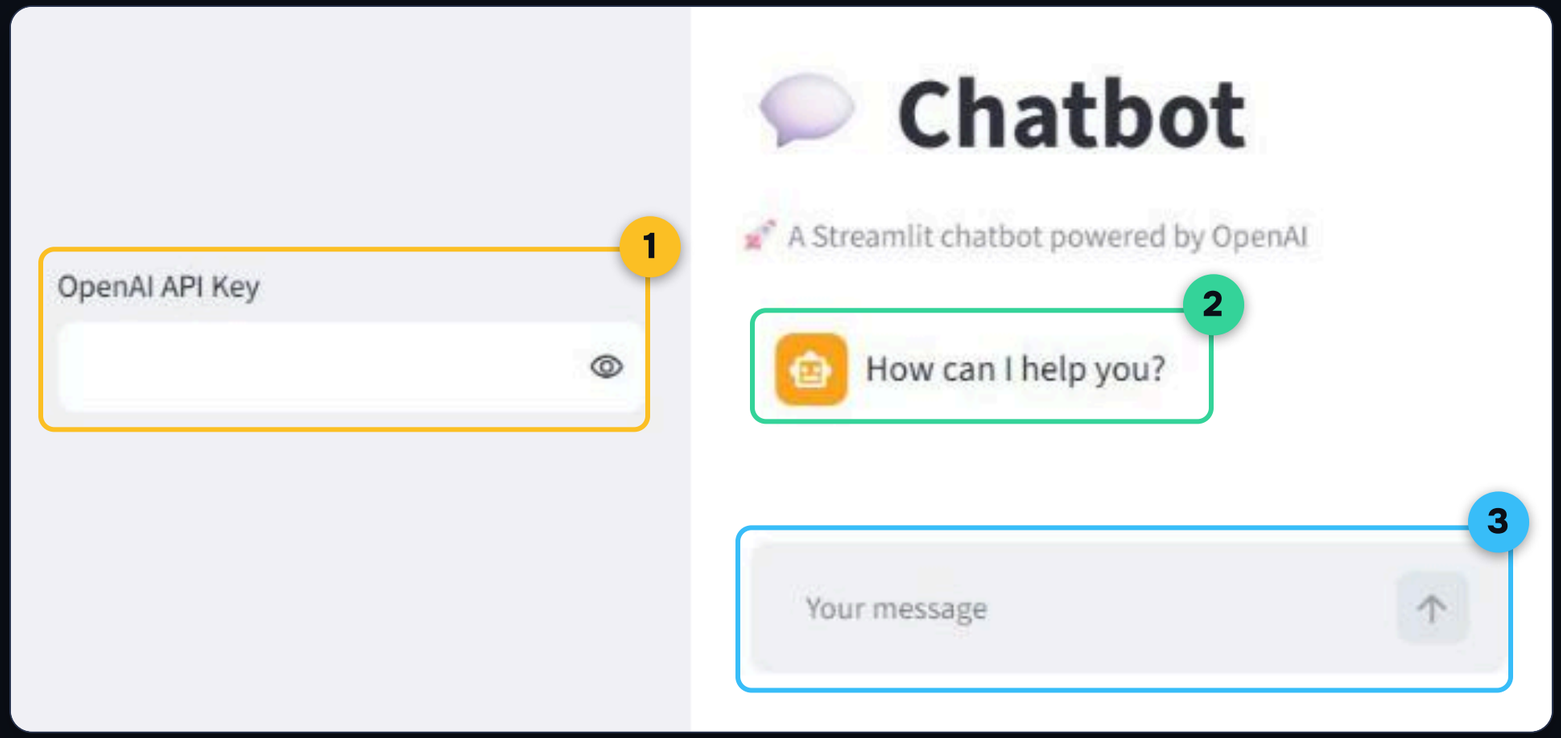


> A clickable prototype gathers requirements better than any meeting: people click, break it and say what they actually miss.

STREAMLIT

Streamlit is a web page in Python: no HTML, no JavaScript

OFFICIAL EXAMPLE LLM-EXAMPLES.STREAMLIT.APP · EXCERPT, CAPTURED 16.10.2026



1. **Sidebar**: a field for the API key
2. **Chat messages**
3. **Input box** at the bottom

> The official example is about 25 lines of Python; our simplified chat is 14 (slide 4). Free, open source (Apache-2.0), version 1.66.

CODE (REFERENCE)

The whole chat is 14 lines: show history, take the question, ask the model, show the answer

APP.PY · RUN: STREAMLIT RUN APP.PY

```
import streamlit as st

def ask_llm(messages):  # stub — replace with your L03 API call
    return "Demo answer: " + messages[-1]["content"]

st.title("Support assistant")

if "messages" not in st.session_state:
    st.session_state.messages = []          # chat memory

for m in st.session_state.messages:          # show history
    st.chat_message(m["role"]).write(m["content"])

if prompt := st.chat_input("Your question"):
    st.session_state.messages.append({"role": "user", "content": prompt})
    st.chat_message("user").write(prompt)
    answer = ask_llm(st.session_state.messages)   # ask the model
    st.session_state.messages.append({"role": "assistant", "content": answer})
    st.chat_message("assistant").write(answer)
```

session_state — chat memory between clicks

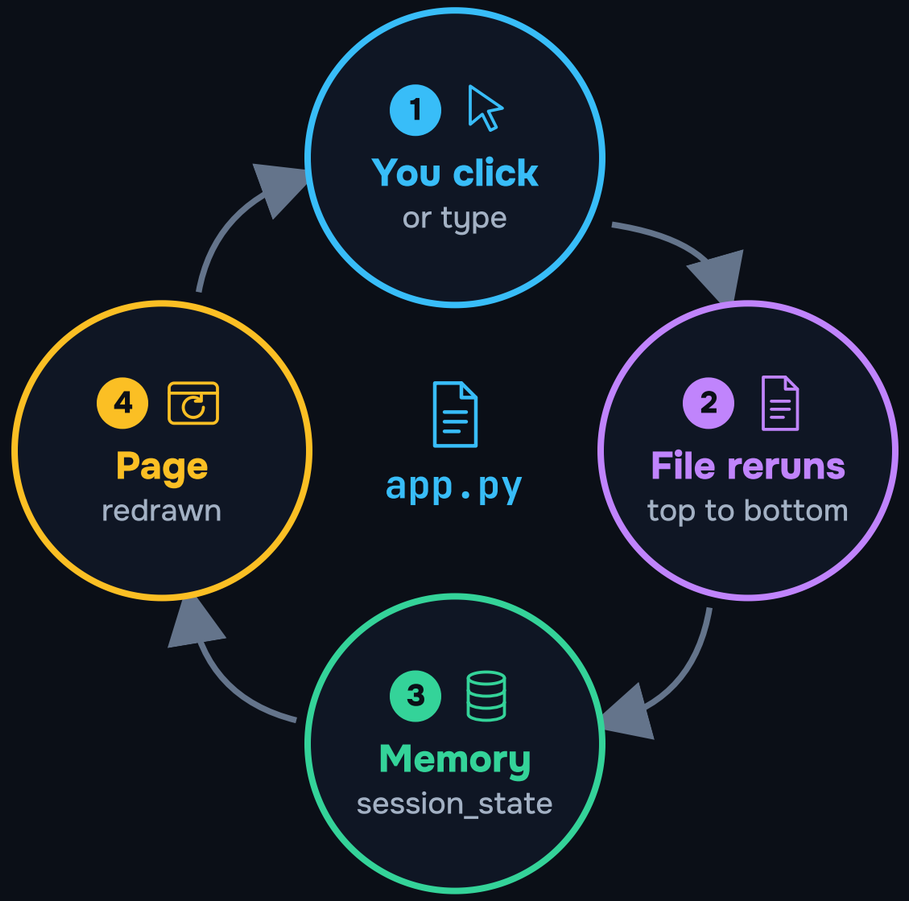
for … chat_message — show the whole history

chat_input — the input box at the bottom

ask_llm — a stub for now; the agent swaps in a model call: Gemini, OpenAI, the local one from L10

HOW IT WORKS

On every click Streamlit reruns the whole file — that is why memory is needed



Anything not stored in **session_state** disappears on the next click.

The most common beginner “bug” — “the chat gets wiped” — is not a bug, it is this loop.

THREE WORDS

Memory, cache and secrets — three things not to mix up

`session_state`

one visitor’s memory: chat history, chosen settings

`cache_data`

a shared cache for everyone: faq.md is read once, not on every click

`secrets`

API keys in .streamlit/secrets.toml, never in code

> A key in code = a key on GitHub (the site where code is stored) = a key for everyone. Add secrets.toml to **.gitignore** — the list of files that never go to GitHub.

WITH AN AGENT

A coding agent writes the code; you set the task and check the result

PROMPT FOR A CODING AGENT: CODEX IN CHATGPT OR ANTIGRAVITY (FREE, WITH LIMITS), CLAUDE CODE (PAID)

```
Build a Streamlit prototype: a chat assistant that answers
staff from faq.md.
1. Python 3.11+, a separate environment:
   pip install streamlit google-genai
2. First run streamlit skills to add Streamlit's official
skills.
3. Model Gemini 3.8 Flash, key from st.secrets; with no key, a
demo mode with a canned answer.
4. History in session_state, read faq.md with st.cache_data.
5. secrets.toml in .gitignore, plus a secrets.example.toml with
a placeholder.
6. Run streamlit run app.py and explain every part of the code
to me.
```

streamlit skills — an official command: it gives the agent instructions on writing Streamlit properly (Claude Code, Codex, Cursor).

A demo mode without a key — the prototype can be shown at once, even without paying for an API.

The last step is a must: ask the agent to explain the code.

RUNNING IT

Four steps: install, run, edit, show

Install

`pip install streamlit`

Python 3.11 or newer

→

Run

`streamlit run app.py`

opens localhost:8501

→

Edit

`save the file → Rerun`

or Always rerun — then automatic

→

Show

`Community Cloud or screen share`

slide 9

> **localhost:8501** — the app runs on your computer. It is not on the internet; on an office network a colleague may open it via the Network URL shown in the terminal.

> Changed text or code — saved — clicked Rerun (the button appears top right). Handy to edit together with a client.

PUBLISHING

Where to show a prototype: your computer, the free Streamlit cloud or your company cloud

Your own computer
free; you, and colleagues on the office network; share your screen to show

Streamlit Community Cloud
free, from a GitHub repo (a code folder); from a public repo it is public; only one private app

Your company
Snowflake (Streamlit's owner), your own cloud — role-based access; engineers decide

A public link = public data and your API bill.
Anyone with the link spends your quota. A public prototype gets synthetic data only.

A free app goes to sleep after 12 hours without visitors.

NO CODE

Sometimes a builder is enough: an app from a description, no code at all

AISTUDIO.GOOGLE.COM → BUILD · CAPTURED 16.10.2026



Google AI Studio Build
describe the app here — get a working one; private by default

Claude Artifacts
a small app right in the chat (L09)

Lovable · Bolt · v0
a whole site from a description; free — a small daily allowance

WHAT TO CHOOSE

Choose by two questions: what you are showing and whether you will read code

	Show an idea	Work with data
I don’t code	**AI Studio Build · Artifacts** in 10 minutes, from a description	**Claude / ChatGPT + files** analysis without an app (L08)
Ready to read code with an agent	**Streamlit** your own UI, your own rules	**Streamlit + pandas** pandas — a table library: charts, your own files

BREAK

5 minutes: think which prototype you will show colleagues

05:00

Next: hands-on, security, the path to a product, homework.

HANDS-ON

10-minute hands-on: a helper prototype for your FAQ

1 Idea
one sentence

→

2 Data
FAQ text, synthetic

→

3 Build
AI Studio Build or an agent

→

4 Test
3 questions

→

5 Show
Teams screen share

No code: AI Studio → Build → “Make a chat helper that answers from this FAQ” + paste 10 Q&As.

With code (only if Python and an agent are already installed): the slide 7 prompt; run streamlit run app.py.

0–7 min — build the prototype and ask it 3 questions. **7–10 min** — 2–3 people share their screen: the question and the prototype’s answer.

SECURITY

Four prototype rules: keys, data, money, sign-in

- **Keys**
  only in secrets, not in code and not in the chat with the agent
- **Data**
  a public prototype — synthetic data only
- **Money**
  a spending limit in the API console; a demo mode with no key
- **Sign-in**
  a private Community Cloud app + email invites

A prototype is not a product. Until engineers and security review it, it does not touch real customer data.

NEXT

What changes when a prototype becomes a product

Prototype	→	Product
one app.py file	→	code in a repo, reviews, tests
synthetic data	→	real data, access rights, logs
a developer’s key	→	a service key, limits, monitoring
works on my machine	→	works for everyone, 24/7

HOMEWORK

Homework: a prototype and feedback from three real people

1. The prototype idea in one sentence: for whom and what it does
2. Build it: Streamlit with an agent, or AI Studio Build
3. Synthetic data only; the key in secrets
4. Show it to 3 colleagues and record their feedback
5. What you would change — 3 points

```
genai-homeworks/L14/
├─ README.md      # idea, conclusion
├─ app.py or a link
├─ faq.md         # synthetic
├─ feedback.md    # 3 people’s feedback
└─ demo.png
```

SUMMARY

A prototype in a day, memory in session_state, keys in secrets, you check the code

A prototype in a day

beats a month of meetings

Streamlit = Python + reruns

memory in session_state

Keys in secrets

a public link = public data

The agent writes, you check

ask it to explain the code

> **Next, L15:** responsible AI, security and law — what you may and may not do with data and content.